


Streams, Sketching and ~~Databases~~ *Big Data*



Graham Cormode

University of Warwick

G.Cormode@Warwick.ac.uk

Big Data

- “Big” data arises in many forms:
 - **Physical Measurements**: from science (physics, astronomy)
 - **Medical data**: genetic sequences, detailed time series
 - **Activity data**: GPS location, social network activity
 - **Business data**: customer behavior tracking at fine detail
- **Common themes**:
 - Data is large, and growing
 - There are important patterns and trends in the data
 - We don’t fully know how to find them

Making sense of Big Data

- Want to be able to interrogate data in different use-cases:
 - **Routine Reporting**: standard set of queries to run
 - **Analysis**: ad hoc querying to answer 'data science' questions
 - **Monitoring**: identify when current behavior differs from old
 - **Mining**: extract new knowledge and patterns from data
- In all cases, need to answer certain basic questions quickly:
 - Describe the distribution of particular attributes in the data
 - How many (distinct) **X** were seen?
 - How many **X < Y** were seen?
 - Give some representative examples of items in the data


Big Data and Hashing

- “Traditional” hashing: compact storage of data
 - Hash tables proportional to data size
 - Fast, compact, exact storage of data
- Hashing with small probability of collisions: very compact storage
 - Bloom filters (no false negatives, bounded false positives)
 - Faster, compacter, probabilistic storage of data
- Hashing with almost certainty of collisions
 - Sketches (items collide, but the signal is preserved)
 - Fasterer, compacterer, approximate storage of data
 - Enables “small summaries for big data”

Data Models

- We model data as a collection of simple **tuples**
- Problems hard due to scale and dimension of input
- Arrivals only model:


- **Example:** $(x, 3), (y, 2), (x, 2)$ encodes the arrival of 3 copies of item x , 2 copies of y , then 2 copies of x .



- Could represent eg. packets on a network; power usage

- Arrivals and departures:


- **Example:** $(x, 3), (y, 2), (x, -2)$ encodes final state of $(x, 1), (y, 2)$.



- Can represent fluctuating quantities, or measure differences between two distributions


Sketches and Frequency Moments

- Sketches as hash-based linear transforms of data
- Frequency distributions and Concentration bounds
- Count-Min sketch for F_∞ and frequent items
- AMS Sketch for F_2
- Estimating F_0
- Extensions:
 - Higher frequency moments
 - Combined frequency moments




Sketch Structures

- **Sketch** is a class of summary that is a **linear transform** of input
 - $\text{Sketch}(x) = Sx$ for some matrix S
 - Hence, $\text{Sketch}(\alpha x + \beta y) = \alpha \text{Sketch}(x) + \beta \text{Sketch}(y)$
 - Trivial to **update** and **merge**
- Often describe S in terms of hash functions
- Aim for limited independence hash functions $h: [n] \rightarrow [m]$
 - If $\Pr_{h \in H} [h(i_1)=j_1 \wedge h(i_2)=j_2 \wedge \dots h(i_k)=j_k] = m^{-k}$,
then H is k -wise independent family (“ h is k -wise independent”)
 - k -wise independent hash functions take time, space $O(k)$



Fingerprints as sketches




- Test if two binary streams are equal
$$d_=(x,y) = 0 \text{ iff } x=y, 1 \text{ otherwise}$$
- To test in small space: pick a suitable hash function h
- Test $h(x)=h(y)$: small chance of false positive, no chance of false negative
- Compute $h(x)$, $h(y)$ incrementally as new bits arrive
 - How to choose the function $h()$?

Polynomial Fingerprints

- Pick $h(x) = \sum_{i=1}^n x_i r^i \bmod p$ for prime p , random $r \in \{1 \dots p-1\}$
- Why?
- Flexible: $h(x)$ is linear function of x —easy to **update** and **merge**
- For accuracy, note that computation **mod** p is over the field \mathbb{Z}_p
 - Consider the polynomial in α , $\sum_{i=1}^n (x_i - y_i) \alpha^i = 0$
 - Polynomial of degree n over \mathbb{Z}_p has at most n roots
- Probability that r happens to solve this polynomial is n/p
- So $\Pr[h(x) = h(y) \mid x \neq y] \leq n/p$
 - Pick $p = \text{poly}(n)$, fingerprints are $\log p = O(\log n)$ bits
- Fingerprints applied to small subsets of data to test equality
 - Will see several examples that use fingerprints as subroutine

Sketches and Frequency Moments

- Sketches as hash-based linear transforms of data
- Frequency distributions and Concentration bounds
- Count-Min sketch for F_∞ and frequent items
- AMS Sketch for F_2
- Estimating F_0
- Extensions:
 - Higher frequency moments
 - Combined frequency moments



Frequency Distributions


- Given set of items, let f_i be the number of occurrences of item i
- Many natural questions on f_i values:
 - Find those i 's with large f_i values (heavy hitters)
 - Find the number of non-zero f_i values (count distinct)
 - Compute $F_k = \sum_i (f_i)^k$ – the k 'th Frequency Moment
 - Compute $H = \sum_i (f_i/F_1) \log (F_1/f_i)$ – the (empirical) entropy
- “Space Complexity of the Frequency Moments”
Alon, Matias, Szegedy in STOC 1996
 - Awarded Gödel prize in 2005
 - Set the pattern for many streaming algorithms to follow

Concentration Bounds

- Will provide randomized algorithms for these problems
- Each algorithm gives a (randomized) estimate of the answer
- Give confidence bounds on the final estimate X
 - Use probabilistic concentration bounds on random variables
- A concentration bound is typically of the form


$$\Pr[|X - x| > \epsilon y] < \delta$$

- At most probability δ of being more than ϵy away from x




Markov Inequality

- Take *any* probability distribution X s.t. $\Pr[X < 0] = 0$
- Consider the event $X \geq k$ for some constant $k > 0$
- For any draw of X , $kI(X \geq k) \leq X$
 - Either $0 \leq X < k$, so $I(X \geq k) = 0$
 - Or $X \geq k$, lhs = k
- Take expectations of both sides: $k \Pr[X \geq k] \leq E[X]$
- **Markov inequality:** $\Pr[X \geq k] \leq E[X]/k$
 - Prob of random variable exceeding k times its expectation $< 1/k$
 - Relatively weak in this form, but still useful




Sketches and Frequency Moments

- Sketches as hash-based linear transforms of data
- Frequency distributions and Concentration bounds
- Count-Min sketch for F_∞ and frequent items
- AMS Sketch for F_2
- Estimating F_0
- Extensions:
 - Higher frequency moments
 - Combined frequency moments




Count-Min Sketch

- Simple **sketch** idea relies primarily on Markov inequality
- Model input data as a vector x of dimension U
- Creates a small summary as an array of $w \times d$ in size
- Use d hash function to map vector entries to $[1..w]$
- Works on arrivals only and arrivals & departures streams



Count-Min Sketch Structure



- Each entry in vector x is mapped to one bucket per row.
- Merge two sketches by entry-wise summation
- Estimate $x[j]$ by taking $\min_k CM[k, h_k(j)]$
 - Guarantees error less than ϵF_1 in size $O(1/\epsilon \log 1/\delta)$
 - Probability of more error is less than $1-\delta$

[C, Muthukrishnan '04]

Approximation of Point Queries

Approximate point query $x'[j] = \min_k CM[k, h_k(j)]$


- Analysis: In k 'th row, $CM[k, h_k(j)] = x[j] + X_{k,j}$
 - $X_{k,j} = \sum_i x[i] I(h_k(i) = h_k(j))$
 - $E[X_{k,j}] = \sum_{i \neq j} x[i] \cdot \Pr[h_k(i) = h_k(j)]$
 $\leq \Pr[h_k(i) = h_k(j)] * \sum_i x[i]$
 $= \varepsilon F_1 / 2$ – requires only pairwise independence of h
 - $\Pr[X_{k,j} \geq \varepsilon F_1] = \Pr[X_{k,j} \geq 2E[X_{k,j}]] \leq 1/2$ by Markov inequality
- So, $\Pr[x'[j] \geq x[j] + \varepsilon F_1] = \Pr[\forall k. X_{k,j} > \varepsilon F_1] \leq 1/2^{\log 1/\delta} = \delta$
- **Final result:** with certainty $x[j] \leq x'[j]$ and with probability at least $1-\delta$, $x'[j] < x[j] + \varepsilon F_1$

Applications of Count-Min to Heavy Hitters

- Count-Min sketch lets us estimate f_i for any i (up to ϵF_1)
- **Heavy Hitters** asks to find i such that f_i is large ($> \phi F_1$)
- **Slow way**: test every i after creating sketch
- **Alternate way**:
 - Keep binary tree over input domain: each node is a subset
 - Keep sketches of all nodes at same level
 - Descend tree to find large frequencies, discard ‘light’ branches
 - Same structure estimates arbitrary range sums
- A first step towards compressed sensing style results...


Application to Large Scale Machine Learning

- In machine learning, often have very large feature space
 - Many objects, each with huge, sparse feature vectors
 - Slow and costly to work in the full feature space
- “Hash kernels”: work with a sketch of the features
 - Effective in practice! [Weinberger, Dasgupta, Langford, Smola, Attenberg '09]
- Similar analysis explains *why*:
 - Essentially, not too much noise on the important features
 - See John Langford’s talk...



Sketches and Frequency Moments

- Frequency distributions and Concentration bounds
- Count-Min sketch for F_∞ and frequent items
- AMS Sketch for F_2
- Estimating F_0
- Extensions:
 - Higher frequency moments
 - Combined frequency moments




Chebyshev Inequality


- Markov inequality applied directly is often quite weak
- But Markov inequality holds for any random variable
- Can apply to a random variable that is a function of X
- Set $Y = (X - E[X])^2$
- By Markov, $\Pr[Y > kE[Y]] < 1/k$
 - $E[Y] = E[(X-E[X])^2] = \text{Var}[X]$
- Hence, $\Pr[|X - E[X]| > \sqrt{k \text{Var}[X]}] < 1/k$
- **Chebyshev inequality:** $\Pr[|X - E[X]| > k] < \text{Var}[X]/k^2$
 - If $\text{Var}[X] \leq \varepsilon^2 E[X]^2$, then $\Pr[|X - E[X]| > \varepsilon E[X]] = O(1)$

F_2 estimation

- AMS sketch (for Alon-Matias-Szegedy) proposed in 1996
 - Allows estimation of F_2 (second frequency moment)
 - Used at the heart of many streaming and non-streaming applications: achieves dimensionality reduction
- Here, describe AMS sketch by generalizing CM sketch.
- Uses extra hash functions $g_1 \dots g_{\log 1/\delta} \{1 \dots U\} \rightarrow \{+1, -1\}$
 - (Low independence) Rademacher variables
- Now, given update $(j, +c)$, set $CM[k, h_k(j)] += c * g_k(j)$



F₂ analysis



- Estimate $F_2 = \text{median}_k \sum_i \text{CM}[k, i]^2$
- Each row's result is $\sum_i g(i)^2 x[i]^2 + \sum_{h(i)=h(j)} 2 g(i) g(j) x[i] x[j]$
- But $g(i)^2 = -1^2 = +1^2 = 1$, and $\sum_i x[i]^2 = F_2$
- $g(i)g(j)$ has 1/2 chance of +1 or -1 : expectation is 0 ...

F₂ Variance

- Expectation of row estimate $R_k = \sum_i CM[k,i]^2$ is exactly F_2
- Variance of row k , $\text{Var}[R_k]$, is an expectation:
 - $\text{Var}[R_k] = E[(\sum_{\text{buckets } b} (CM[k,b])^2 - F_2)^2]$
 - Good exercise in algebra: expand this sum and simplify
 - Many terms are zero in expectation because of terms like $g(a)g(b)g(c)g(d)$ (degree at most 4)
 - Requires that hash function g is *four-wise independent*: it behaves uniformly over subsets of size four or smaller
 - Such hash functions are easy to construct

F₂ Variance

- Terms with odd powers of $g(a)$ are zero in expectation
 - $g(a)g(b)g^2(c), g(a)g(b)g(c)g(d), g(a)g^3(b)$

- Leaves

$$\begin{aligned}\text{Var}[R_k] &\leq \sum_i g^4(i) x[i]^4 \\ &\quad + 2 \sum_{j \neq i} g^2(i) g^2(j) x[i]^2 x[j]^2 \\ &\quad + 4 \sum_{h(i)=h(j)} g^2(i) g^2(j) x[i]^2 x[j]^2 \\ &\quad - (x[i]^4 + \sum_{j \neq i} 2x[i]^2 x[j]^2) \\ &\leq F_2^2/w\end{aligned}$$


- Row variance can finally be bounded by F_2^2/w
 - Chebyshev for $w=4/\varepsilon^2$ gives probability $\frac{1}{4}$ of failure:
 $\Pr[|R_k - F_2| > \varepsilon^2 F_2] \leq \frac{1}{4}$
 - How to amplify this to small δ probability of failure?
 - Rescaling w has cost linear in $1/\delta$

Tail Inequalities for Sums

- We achieve stronger bounds on tail probabilities for the sum of independent *Bernoulli trials* via the **Chernoff Bound**:
 - Let X_1, \dots, X_m be **independent** Bernoulli trials s.t. $\Pr[X_i=1] = p$ ($\Pr[X_i=0] = 1-p$).
 - Let $X = \sum_{i=1}^m X_i$, and $\mu = mp$ be the expectation of X .
 - $\Pr[X > (1+\varepsilon)\mu] = \Pr[\exp(tX) > \exp(t(1+\varepsilon)\mu)] \leq E[\exp(tX)]/\exp(t(1+\varepsilon)\mu)$
 - $E[\exp(tX)] = \prod_i E[\exp(tX_i)] = \prod_i (1-p + pe^t) \leq \prod_i \exp(p(e^t-1))$
 $= \exp(\mu(e^t - 1))$
 - $\Pr[X > (1+\varepsilon)\mu] \leq \exp(\mu(e^t - 1) - \mu t(1+\varepsilon)) = \exp(\mu(-\varepsilon t + t^2/2 + t^3/6 + \dots))$
 $\leq \exp(\mu(t^2/2 - \varepsilon t))$
 - **Balance**: choose $t=\varepsilon/2$ $\leq \exp(-\mu \varepsilon^2/2)$

Applying Chernoff Bound

- Each row gives an estimate that is within ε relative error with probability $p' > 3/4$
- Take d repetitions and find the median. Why the median?




- Because bad estimates are either too small or too large
- Good estimates form a contiguous group “in the middle”
- At least $d/2$ estimates must be bad for median to be bad
- Apply Chernoff bound to d independent estimates, $p=1/4$
 - $\Pr[\text{More than } d/2 \text{ bad estimates}] < 2\exp(-d/8)$
 - So we set $d = \Theta(\ln 1/\delta)$ to give δ probability of failure
- Same outline used many times in summary construction

Applications and Extensions

- F_2 guarantee: estimate $\|x\|_2$ from sketch with error $\varepsilon \|x\|_2$
 - Since $\|x + y\|_2^2 = \|x\|_2^2 + \|y\|_2^2 + 2x \cdot y$
Can estimate $(x \cdot y)$ with error $\varepsilon \|x\|_2 \|y\|_2$
 - If $y = e_j$, obtain $(x \cdot e_j) = x_j$ with error $\varepsilon \|x\|_2$:
 L_2 guarantee (“Count Sketch”) vs L_1 guarantee (Count-Min)
- Can view the sketch as a low-independence realization of the Johnson-Lindendestraus lemma
 - Best current JL methods have the same structure
 - **JL is stronger**: embeds directly into Euclidean space
 - **JL is also weaker**: requires $O(1/\varepsilon)$ -wise hashing, $O(\log 1/\delta)$ independence [Nelson, Nguyen 13]

Sketches and Frequency Moments

- Frequency Moments and Sketches
- Count-Min sketch for F_∞ and frequent items
- AMS Sketch for F_2
- Estimating F_0
- Extensions:
 - Higher frequency moments
 - Combined frequency moments




F_0 Estimation

- F_0 is the number of distinct items in the stream
 - a fundamental quantity with many applications
- Early algorithms by [Flajolet and Martin \[1983\]](#) gave nice hashing-based solution
 - analysis assumed fully independent hash functions
- Will describe a generalized version of the FM algorithm due to [Bar-Yossef et. al](#) with only pairwise indendence
 - Known as the “k-Minimum values (KMV)” algorithm

F_0 Algorithm


- Let m be the domain of stream elements
 - Each item in data is from $[1\dots m]$
- Pick a random (pairwise) hash function $h: [m] \rightarrow [R]$
 - For $R = m^3$ with probability at least $1-1/m$, no collisions under h



- For each stream item i , compute $h(i)$, and track the t distinct items achieving the smallest values of $h(i)$
 - **Note:** if same i is seen many times, $h(i)$ is same
 - Let $v_t = t$ 'th smallest (distinct) value of $h(i)$ seen
- If $n = F_0 < t$, give exact answer, else estimate $F'_0 = tR/v_t$
 - $v_t/R \approx$ fraction of hash domain occupied by t smallest

Analysis of F_0 algorithm

- Suppose $F'_0 = tR/v_t > (1+\varepsilon)n$ [estimate is too high]



- So for input = set $S \in 2^{[m]}$, we have
 - $|\{s \in S \mid h(s) < tR/(1+\varepsilon)n\}| > t$
 - Because $\varepsilon < 1$, we have $tR/(1+\varepsilon)n \leq (1-\varepsilon/2)tR/n$
 - $\Pr[h(s) < (1-\varepsilon/2)tR/n] \approx 1/R * (1-\varepsilon/2)tR/n = (1-\varepsilon/2)t/n$
 - (this analysis outline hides some rounding issues)

Chebyshev Analysis

- Let Y be number of items hashing to under $tR/(1+\epsilon)n$
 - $E[Y] = n * \Pr[h(s) < tR/(1+\epsilon)n] = (1-\epsilon/2)t$
 - For each item i , variance of the event = $p(1-p) < p$
 - $\text{Var}[Y] = \sum_{s \in S} \text{Var}[h(s) < tR/(1+\epsilon)n] < (1-\epsilon/2)t$
 - We sum variances because of pairwise independence
- Now apply **Chebyshev inequality**:
 - $\Pr[Y > t] \leq \Pr[|Y - E[Y]| > \epsilon t/2]$
$$\leq 4\text{Var}[Y]/\epsilon^2 t^2$$
$$< 4t/(\epsilon^2 t^2)$$
 - Set $t=20/\epsilon^2$ to make this $\text{Prob} \leq 1/5$

Completing the analysis

- We have shown

$$\Pr[F'_0 > (1+\varepsilon) F_0] < 1/5$$

- Can show $\Pr[F'_0 < (1-\varepsilon) F_0] < 1/5$ similarly

- too few items hash below a certain value

- So $\Pr[(1-\varepsilon) F_0 \leq F'_0 \leq (1+\varepsilon) F_0] > 3/5$ [Good estimate]

- Amplify this probability: repeat $O(\log 1/\delta)$ times in parallel with different choices of hash function h

- Take the median of the estimates, analysis as before

F_0 Issues

■ Space cost:

- Store t hash values, so $O(1/\varepsilon^2 \log m)$ bits
- Can improve to $O(1/\varepsilon^2 + \log m)$ with additional tricks



■ Time cost:

- Find if hash value $h(i) < v_t$
- Update v_t and list of t smallest if $h(i)$ not already present
- Total time $O(\log 1/\varepsilon + \log m)$ worst case

Count-Distinct

- Engineering the best constants: **Hyperloglog algorithm**
 - Hash each item to one of $1/\epsilon^2$ buckets (like Count-Min)
 - In each bucket, track the function $\max \lfloor \log(h(x)) \rfloor$
 - Can view as a coarsened version of KMV
 - Space efficient: need $\log \log m \approx 6$ bits per bucket
- Can estimate intersections between sketches
 - Make use of identity $|A \cap B| = |A| + |B| - |A \cup B|$
 - Error scales with $\epsilon \sqrt{|A| |B|}$, so poor for small intersections
 - Higher order intersections via inclusion-exclusion principle

Subset Size Estimation from KMV

- Want to estimate the fraction $f = |A|/|S|$
 - S is the observed set of data
 - A is an arbitrary subset given later
 - E.g. fraction of customers who are female 18-24 from Denmark
- Simple algorithm:
 - Run KMV to get sample set K , estimate $f' = |A \cap K|/k$
 - Need to bound probability of getting a bad estimate
 - Analysis due to [Thorup 13]

Subset Size Estimation

- Upper bound:
 - Suppose we overestimate: $|A \cap K| > (1 + a) / (1 - b) fk$
 - Set threshold $t = kR/(n(1-a))$
- To have overestimate, must have one of:
 1. Fewer than k elements from B hash below t : expect $k/(1-a)$
 2. More than $(1+b)(kf)/(1-a)$ elements from A hash below t :
expect $kf/(1-a)$
 - Otherwise, cannot have overestimate
- To analyze, bound the probability of 1. and 2. separately
 - Probability of overestimate is bounded by sum of these probs

Bounding error probability


- Use Chebyshev to bound the two bad cases
 - Suppose mean number of m hash values below a threshold $\mu = mp$
 - Standard deviation $\sigma = ((1-p)pm)^{1/2} \leq \mu^{1/2}$ (via pairwise independence)
 - Set $a = 4/\sqrt{k}$, $b = 4/\sqrt{fk}$
 - For Event 1., we have $\mu = k/(1-a) \geq k$ so, via Chebyshev,
 $\Pr[\text{Event 1.}] \leq \mu/a\sigma < 1/16$
 - Similarly, for Event 2., we have $\mu = kf/(1-a) \geq kf$ so
 $\Pr[\text{Event 2.}] \leq \mu/b\sigma < 1/16$
 - By union bound, at most $1/8$ prob of overestimate
- Similar case analysis for the case of an underestimate

Subset count accuracy

- With probability at least $\frac{3}{4}$, the error is $O((fk)^{\frac{1}{2}})$
 - Arises from the choice of parameters b and a
 - Error scales with f
- For some lower bound on f , f' , can get relative error ε :
 - Set $k \propto f'/\varepsilon^2$ for $(1 \pm \varepsilon)$ error with constant probability
- For improved error:
 - Either increase $k \propto 1/\delta$
 - Or repeat $\log 1/\delta$ times and take median estimate

Frequency Moments

- Intro to frequency distributions and Concentration bounds
- Count-Min sketch for F_∞ and frequent items
- AMS Sketch for F_2
- Estimating F_0
- **Extensions:**
 - Higher frequency moments
 - Combined frequency moments



Higher Frequency Moments

- F_k for $k > 2$. Use a sampling trick [Alon et al 96]:
 - Uniformly pick an item from the stream length $1 \dots n$
 - Set r = how many times that item appears subsequently
 - Set estimate $F'_k = n(r^k - (r-1)^k)$
- $E[F'_k] = 1/n * n * [f_1^k - (f_1-1)^k + (f_1-1)^k - (f_1-2)^k + \dots + 1^k - 0^k] + \dots$
 $= f_1^k + f_2^k + \dots = F_k$
- $\text{Var}[F'_k] \leq 1/n * n^2 * [(f_1^k - (f_1-1)^k)^2 + \dots]$
 - Use various bounds to bound the variance by $k m^{1-1/k} F_k^2$
 - Repeat $k m^{1-1/k}$ times in parallel to reduce variance
- Total space needed is $O(k m^{1-1/k})$ machine words
 - Not a sketch: does not distribute easily. See next lecture!

Combined Frequency Moments

- Let $G[i,j] = 1$ if (i,j) appears in input.
E.g. graph edge from i to j . Total of m distinct edges
- Let $d_i = \sum_{j=1}^n G[i,j]$ (aka degree of node i)
- Find aggregates of d_i 's:
 - Estimate heavy d_i 's (people who talk to many)
 - Estimate frequency moments:
number of distinct d_i values, sum of squares
 - Range sums of d_i 's (subnet traffic)
- **Approach**: nest one sketch inside another, e.g. HLL inside CM
 - Requires new analysis to track overall error

Range Efficiency

- Sometimes input is specified as a collection of ranges $[a,b]$
 - $[a,b]$ means insert all items $(a, a+1, a+2 \dots b)$
 - Trivial solution: just insert each item in the range
- Range efficient F_0 [Pavan, Tirthapura 05]
 - Start with an alg for F_0 based on pairwise hash functions
 - Key problem: track which items hash into a certain range
 - Dives into hash fns to divide and conquer for ranges
- Range efficient F_2 [Calderbank et al. 05, Rusu,Dobra 06]
 - Start with sketches for F_2 which sum hash values
 - Design new hash functions so that range sums are fast
- Rectangle Efficient F_0 [Tirthapura, Woodruff 12]

Summary

- Sketching Techniques summarize large data sets
- Summarize vectors:
 - Test equality (fingerprints)
 - Recover approximate entries (count-min, count sketch)
 - Approximate Euclidean norm (F_2) and dot product
 - Approximate number of non-zero entries (F_0)
 - Approximate set membership (Bloom filter)

Current Directions in Streaming and Sketching

- **Sparse representations** of high dimensional objects
 - Compressed sensing, sparse fast fourier transform
- **Numerical linear algebra** for (large) matrices
 - k-rank approximation, linear regression, PCA, SVD, eigenvalues
- Computations on large **graphs**
 - Sparsification, clustering, matching
- **Geometric** (big) data
 - Coresets, facility location, optimization, machine learning
- Use of summaries in **distributed computation**
 - MapReduce, Continuous Distributed models