

Hashing for Machine Learning

Hashing Summer School, University of Copenhagen

To follow along:

```
git clone git://github.com/JohnLangford/vowpal_wabbit.git  
wget http://hunch.net/~jl/rcv1.tar.gz
```



John Langford, Microsoft Research, NYC

July 17, 2014

Simple Machine Learning: Linear Learning

Features: a vector $x \in \mathbb{R}^n$

Label: $y \in \mathbb{R}$

Goal: **Learn** $w \in \mathbb{R}^n$ such that $\hat{y}_w(x) = \sum_i w_i x_i$ is close to y .

An Example: The RCV1 dataset

Pick whether a document is in category CCAT or not.

Dataset size:

781K examples

60M nonzero features

1.1G bytes

Format: label | sparse features ...

An Example: The RCV1 dataset

Pick whether a document is in category CCAT or not.

Dataset size:

781K examples

60M nonzero features

1.1G bytes

Format: label | sparse features ...

1 | 13:3.9656971e-02 24:3.4781646e-02 ...

An Example: The RCV1 dataset

Pick whether a document is in category CCAT or not.

Dataset size:

781K examples

60M nonzero features

1.1G bytes

Format: **label** | **sparse features ...**

1 | **13:3.9656971e-02 24:3.4781646e-02 ...**

which corresponds to:

1 | **tuesday year ...**

with “Term Frequency Inverse Document Frequency” feature values.

An Example: The RCV1 dataset

Pick whether a document is in category CCAT or not.

Dataset size:

781K examples

60M nonzero features

1.1G bytes

Format: label | sparse features ...

1 | 13:3.9656971e-02 24:3.4781646e-02 ...

which corresponds to:

1 | tuesday year ...

with “Term Frequency Inverse Document Frequency” feature values.

run: `vw rcv1.train.txt -c -binary`

Progressive validation loss 0.0513 (\simeq best) in 1 second. (best)

An Example: The RCV1 dataset

Pick whether a document is in category CCAT or not.

Dataset size:

781K examples

60M nonzero features

1.1G bytes

Format: **label** | **sparse features ...**

1 | **13:3.9656971e-02 24:3.4781646e-02 ...**

which corresponds to:

1 | **tuesday year ...**

with “Term Frequency Inverse Document Frequency” feature values.

run: **vw rcv1.train.txt -c -binary**

Progressive validation loss **0.0513** (\simeq best) in **1 second**. (best)

Performance real, but timing not: TFIDF transform >2 minutes.

The unconventional story: Hashing on RCV1

Start with the raw data:

1 | tuesday year million short compan vehicl line stat financ commit
exchang plan corp subsid credit issu debt pay gold bureau prelimin
refin billion telephon time draw

...

The unconventional story: Hashing on RCV1

Start with the raw data:

```
1 | tuesday year million short compan vehicl line stat financ commit  
exchang plan corp subsid credit issu debt pay gold bureau prelimin  
refin billion telephon time draw
```

...

```
vw -b 24 -c rcv1.train.raw.txt --binary
```

Progressive validation loss 0.0572 in 1.3s

The unconventional story: Hashing on RCV1

Start with the raw data:

1 | tuesday year million short compan vehicl line stat financ commit
exchang plan corp subsid credit issu debt pay gold bureau prelimin
refin billion telephon time draw

...

```
vw -b 24 -c rcv1.train.raw.txt --binary
```

Progressive validation loss 0.0572 in 1.3s

```
vw -b 24 -c rcv1.train.raw.txt --binary --ngram 2
```

Progressive validation loss 0.0500 in 3.4s

The unconventional story: Hashing on RCV1

Start with the raw data:

1 | tuesday year million short compan vehicl line stat financ commit
exchang plan corp subsid credit issu debt pay gold bureau prelimin
refin billion telephon time draw

...

```
vw -b 24 -c rcv1.train.raw.txt --binary
```

Progressive validation loss 0.0572 in 1.3s

```
vw -b 24 -c rcv1.train.raw.txt --binary --ngram 2
```

Progressive validation loss 0.0500 in 3.4s

```
vw -b 24 -c rcv1.train.raw.txt --binary --ngram 2 --skips 4
```

Progressive validation loss 0.0454 in 11s

The unconventional story: Hashing on RCV1

Start with the raw data:

1 | tuesday year million short compan vehicl line stat financ commit
exchang plan corp subsid credit issu debt pay gold bureau prelimin
refin billion telephon time draw

...

```
vw -b 24 -c rcv1.train.raw.txt --binary
```

Progressive validation loss 0.0572 in 1.3s

```
vw -b 24 -c rcv1.train.raw.txt --binary --ngram 2
```

Progressive validation loss 0.0500 in 3.4s

```
vw -b 24 -c rcv1.train.raw.txt --binary --ngram 2 --skips 4
```

Progressive validation loss 0.0454 in 11s

```
vw -b 24 -c rcv1.train.raw.txt --binary --ngram 2 --skips 4 -l 0.25
```

Progressive validation loss 0.0449 in 11s

+Better error rate

+1/10th traing time!

+faster/easier testing

Outline

- ① A bit more about machine Learning
- ② Use of Hash
- ③ Hashing applications

Algorithm used: Online Linear Learning

Start with $\forall i : w_i = 0$

Repeatedly:

- 1 Get features $x \in \mathbb{R}^n$.
- 2 Make linear prediction $\hat{y}_w(x) = \sum_i w_i x_i$.
- 3 Observe label y .
- 4 Record loss $\ell(y, \hat{y}_w(x))$.
- 5 Update weights so $\hat{y}_w(x)$ is closer to y .

Algorithm used: Online Linear Learning

Start with $\forall i : w_i = 0$

Repeatedly:

- 1 Get **features** $x \in \mathbb{R}^n$.
- 2 Make linear **prediction** $\hat{y}_w(x) = \sum_i w_i x_i$.
- 3 Observe **label** y .
- 4 Record **loss** $\ell(y, \hat{y}_w(x))$.
- 5 **Update** weights so $\hat{y}_w(x)$ is closer to y .
Example: $w_i \leftarrow w_i + \eta(y - \hat{y})x_i$.

Algorithm used: Online Linear Learning

Start with $\forall i : w_i = 0$

Repeatedly:

- 1 Get **features** $x \in \mathbb{R}^n$.
- 2 Make linear **prediction** $\hat{y}_w(x) = \sum_i w_i x_i$.
- 3 Observe **label** y .
- 4 Record **loss** $\ell(y, \hat{y}_w(x))$.
- 5 **Update** weights so $\hat{y}_w(x)$ is closer to y .
Example: $w_i \leftarrow w_i + \eta(y - \hat{y})x_i$.

(many more details that matter in practice)

Features in Practice: Engineered Features

Hand crafted features, built up iteratively over time, each new feature fixing a discovered problem.

- ① +Good understanding of what's happening.
- ② +Never fail to learn the obvious.
- ③ +Small RAM usage.
- ④ -Slow at test time. Intuitive features for humans can be hard
- ⑤ -Low Capacity. A poor fit for large datasets. (Boosted)
Decision trees are a good compensation on smaller datasets.
- ⑥ -High persontime.

Features in Practice: Learned Features

Use a nonlinear/nonconvex possibly deep learning algorithm.

- ① +Good results in Speech & Vision.
- ② +Fast at test time.
- ③ +High capacity. Useful on large datasets.
- ④ -Slow training. Days to weeks are common.
- ⑤ -Wizardry may be required.

Features in Practice: sparse words

Generate a feature for every word.

- ① +High capacity.
- ② +fast test time. Lookup some numbers, then compute an easy prediction.
- ③ +fast learning Linear faster than decision tree, but parallel is tricky.
- ④ -High RAM usage

Features in Practice: sparse words

Generate a feature for every word.

- ① +High capacity.
- ② +fast test time. Lookup some numbers, then compute an easy prediction. This lecture.
- ③ +fast learning Linear faster than decision tree, but parallel is tricky. This lecture.
- ④ -High RAM usage This lecture.

Outline

- 1 A bit more about machine Learning
- 2 Use of Hash
- 3 Hashing applications

What is hashing?

Hash function: $\text{string} \rightarrow \{0, 1\}^b$

What is hashing?

Hash function: $\text{string} \rightarrow \{0, 1\}^b$

Hash table = Hash function + Array< Pair<string, int> > of
length $\{0, 1\}^b$

What is hashing?


Hash function: $\text{string} \rightarrow \{0, 1\}^b$

Hash table = Hash function + Array< Pair<string, int> > of length $\{0, 1\}^b$

Perfect hash = mapping of n fixed (and known in advance) strings to integers $\{1, n\}$.

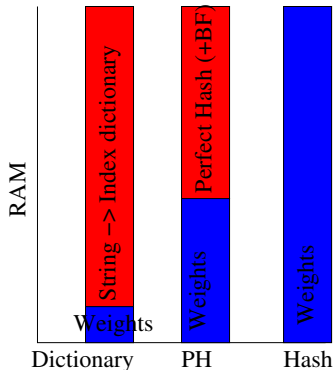
How does feature address parameter?

- 1 **Hash Table** (aka Dictionary): Store hash function + Every string + Index.
- 2 **Perfect Hash** (+Bloom Filter): Store Custom Hash function (+ bit array).
- 3 **Hash function**: Store Hash function.



How does feature address parameter?

- 1 **Hash Table** (aka Dictionary): Store hash function + Every string + Index.
- 2 **Perfect Hash** (+Bloom Filter): Store Custom Hash function (+ bit array).
- 3 **Hash function**: Store Hash function.



More weights is better!

A variation: unbiased hashing

Multiply feature value by $(-1)^s$ where s is a 1 bit hash.

Advantage: Feature values have expectation 0 for a random $s \Rightarrow$ better performance in the high collision regime.

Objection: Collisions!

Valid sometimes: particularly with low dimensional hand engineered features.

Objection: Collisions!

Valid sometimes: particularly with low dimensional hand engineered features.

Theorem: If a feature is duplicated $O(\log n)$ times when there are $O(n)$ features, at least one version of the feature is uncollided when hashing with $\log(n \log n)$ bits.

Proof: Essentially Bloom filter logic. See Michael Mitzenmacher's talk Monday.

Objection: Incomprehensible!

Objection: Incomprehensible!

Use `--audit` to decode.

Keep your own dictionary on the side `--invert_hash` if needed.

Outline

- 1 A bit more about machine Learning
- 2 Use of Hash
- 3 Hashing applications

Use of Hash: Ngrams

2-gram = a feature for every pair of adjacent words.

3-gram = a feature for every triple of adjacent words, etc...

n-gram = ...

Use of Hash: Ngrams

2-gram = a feature for every pair of adjacent words.

3-gram = a feature for every triple of adjacent words, etc...

n-gram = ...

Input:

$(\text{index}_1, \text{value}_1)$

$(\text{index}_2, \text{value}_2)$

Output:

Use of Hash: Ngrams

2-gram = a feature for every pair of adjacent words.

3-gram = a feature for every triple of adjacent words, etc...

n-gram = ...

Input:

$(\text{index}_1, \text{value}_1)$

$(\text{index}_2, \text{value}_2)$

Output:

$((\text{index}_1 \text{magic} + \text{index}_2) \& \text{mask}, \text{value}_1 \text{value}_2)$

(linear hash, value multiplication)

Use of Hash: Outer Products

Input:

Feature sets F_1, F_2

Output:

Outer product set $F_1 \times F_2$

Use linear hash again.

Use of Hash: Outer Products

Input:

Feature sets F_1, F_2

Output:

Outer product set $F_1 \times F_2$

Use linear hash again.

Implemented via $-q$ in VW.

Example 2: Mass Personalized Spam Filtering

- 1 $3.2 * 10^6$ labeled emails.
- 2 433167 users.
- 3 $\sim 40 * 10^6$ unique tokens.

How do we construct a spam filter which is personalized, yet uses global information?

Example 2: Mass Personalized Spam Filtering

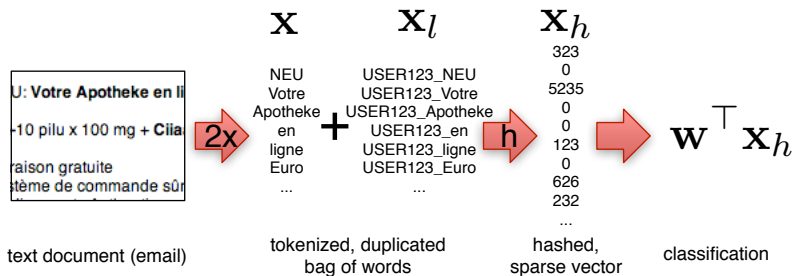
- 1 $3.2 * 10^6$ labeled emails.
- 2 433167 users.
- 3 $\sim 40 * 10^6$ unique tokens.

How do we construct a spam filter which is personalized, yet uses global information?

Bad answer: Construct a global filter + 433167 personalized filters using a conventional hashmap to specify features. This might require $433167 * 40 * 10^6 * 4 \sim 70$ Terabytes of RAM.

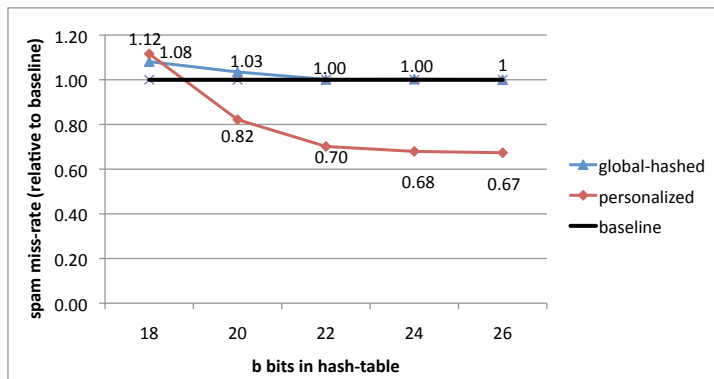
Using Hashing

Use hashing to predict according to: $\langle w, \phi(x) \rangle + \langle w, \phi_u(x) \rangle$



(in VW: specify the userid as a feature and use -q)

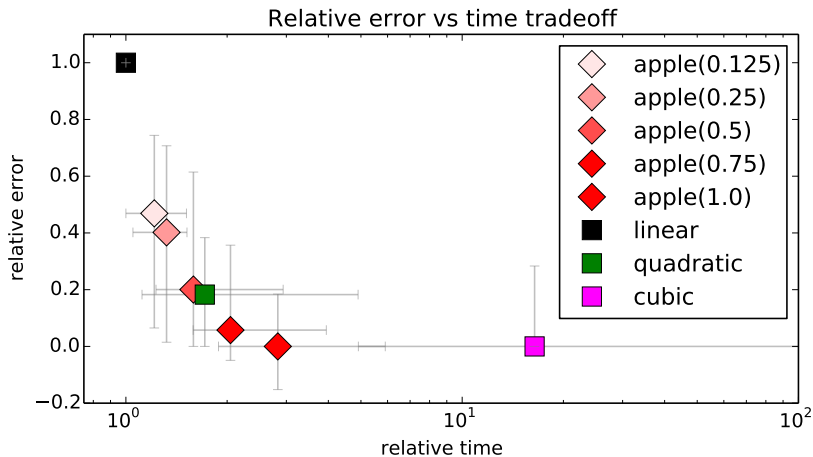
Results



2^{26} parameters = 64M parameters = 256MB of RAM.

An **x270K** savings in RAM requirements.

Another Application: Efficient Sparse Polynomial learning



Features sometimes collide, which is scary, but then you love it

Generate a feature for every word, ngram, skipgram, pair of (ad word, query word), etc... and use high dimensional representation.

- ① +High capacity.
- ② +Correlation effects nailed.
- ③ +Fast test time. Compute an easy prediction.
- ④ +Fast Learning (with Online + parallel techniques. See talks.)
- ⑤ +/-Variable RAM usage. Highly problem dependent but fully controlled.

Another cool observation: Online learning + Hashing = learning algorithm with fully controlled memory footprint \Rightarrow Robustness.

References, prequels

- ① Reinforcement Learning: An Introduction, Richard S. Sutton and Andrew G. Barto, MIT Press, Cambridge, MA, 1998. Chapter 8.3.1 hashes states.
- ② CRM114 <http://crm114.sourceforge.net/>, 2002. Uses hashing of grams for spam detection.
- ③ Apparently used by others as well, internally.
- ④ Many use hashables which store the original item or a 64+ bit hash of the original item.

References, “modern” hashing trick

- 1 2007, Langford, Li, Strehl, Vowpal Wabbit released.
- 2 2008, Ganchev & Dredze, ACL workshop: A hash function is as good as a hashmap empirically.
- 3 2008/2009, VW Reimplementation/Reimagination/Integration in Stream (James Patterson & Alex Smola) and Torch (Jason Weston, Olivier Chapelle, Kilian).
- 4 2009, AISTAT Qinfeng Shi et al, Hash kernel definition, Asymptopia Redundancy analysis
- 5 2009, ICML Kilian et al, Unbiased Hash Kernel, Length Deviation Bound, Mass Personalization Example and Multiuse Bound.

Question 1

Hashing is applied to **features** before learning. Is it better to apply to **parameters** after learning?

Question 2

A commonly proposed alternative to hashing is **random projection** as per Johnson-Lindenstrauss (see Alex Andoni's lecture yesterday). Is that a better approach?

Question 3

How do you compute the n -grams for k elements in time $O(k)$?