

Adding Back-pressure to Spark Streaming

Dean Wampler, Typesafe

Typesafe and Spark

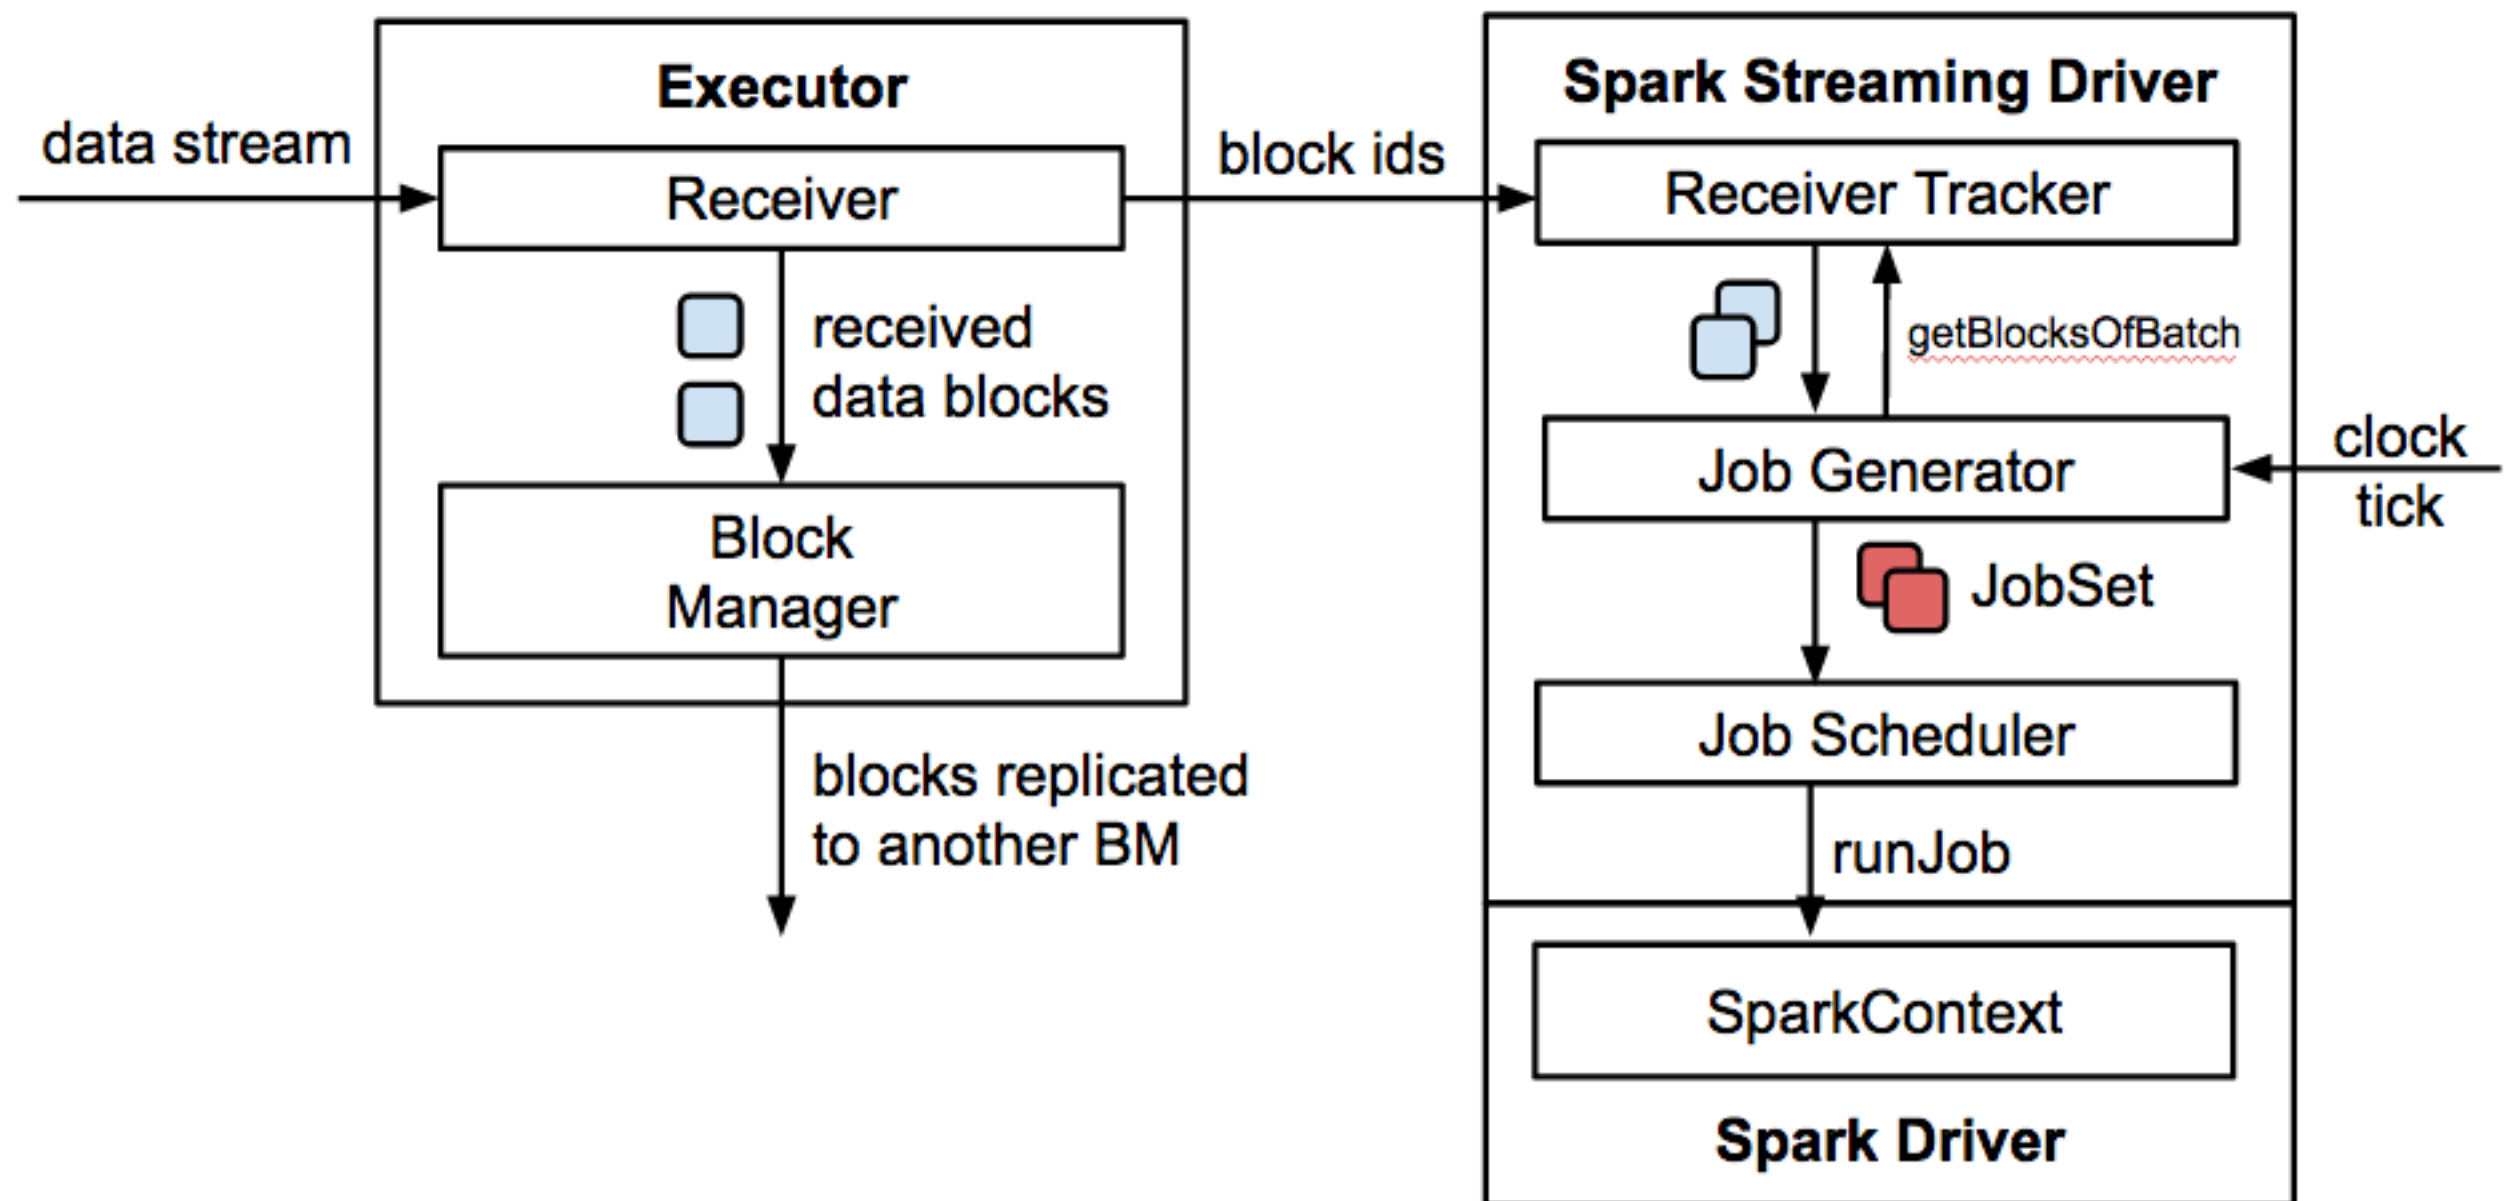
- Spark aligned with our passion for distributed, streaming (“fast data”) systems... and Scala!
- We’re contributing to Spark’s use of Scala, Streaming, and Mesos integration.
- Commercial support for Spark on Mesos.
 - *Mesosphere Infinity.*

The problem


Data may arrive faster than Spark can process it, leading to an unstable system (delays, memory spills).

Worst-case, you run out of memory!

Spark Streaming Architecture



If the batch interval is too short, data accumulates in the system



Rate Limiting in Spark 1.4

- You could set a **static** rate limit
 - Had to be conservative
 - difficult to find the right limit (depends on cluster size)
 - one limit to rule all streams for all time!

Back-pressure



Back-pressure

- (aka, flow control)
- a slow *consumer* should slow down the *producer*
- Classic example, TCP
- New example: Reactive Streams
 - reactive-streams.org

Back-pressure in Spark 1.5

- *A dynamic* rate limiter
 1. Estimate the number of elements that can be safely processed by the system in the *next* batch interval
 2. Communicate the limit to Receivers
 3. Rely on TCP to slow down producers on the other side of the channel

Note: What's Missing?

- What about other distributed streams?
- Do we want to rely solely on TCP back-pressure?
- We'll come back to this...

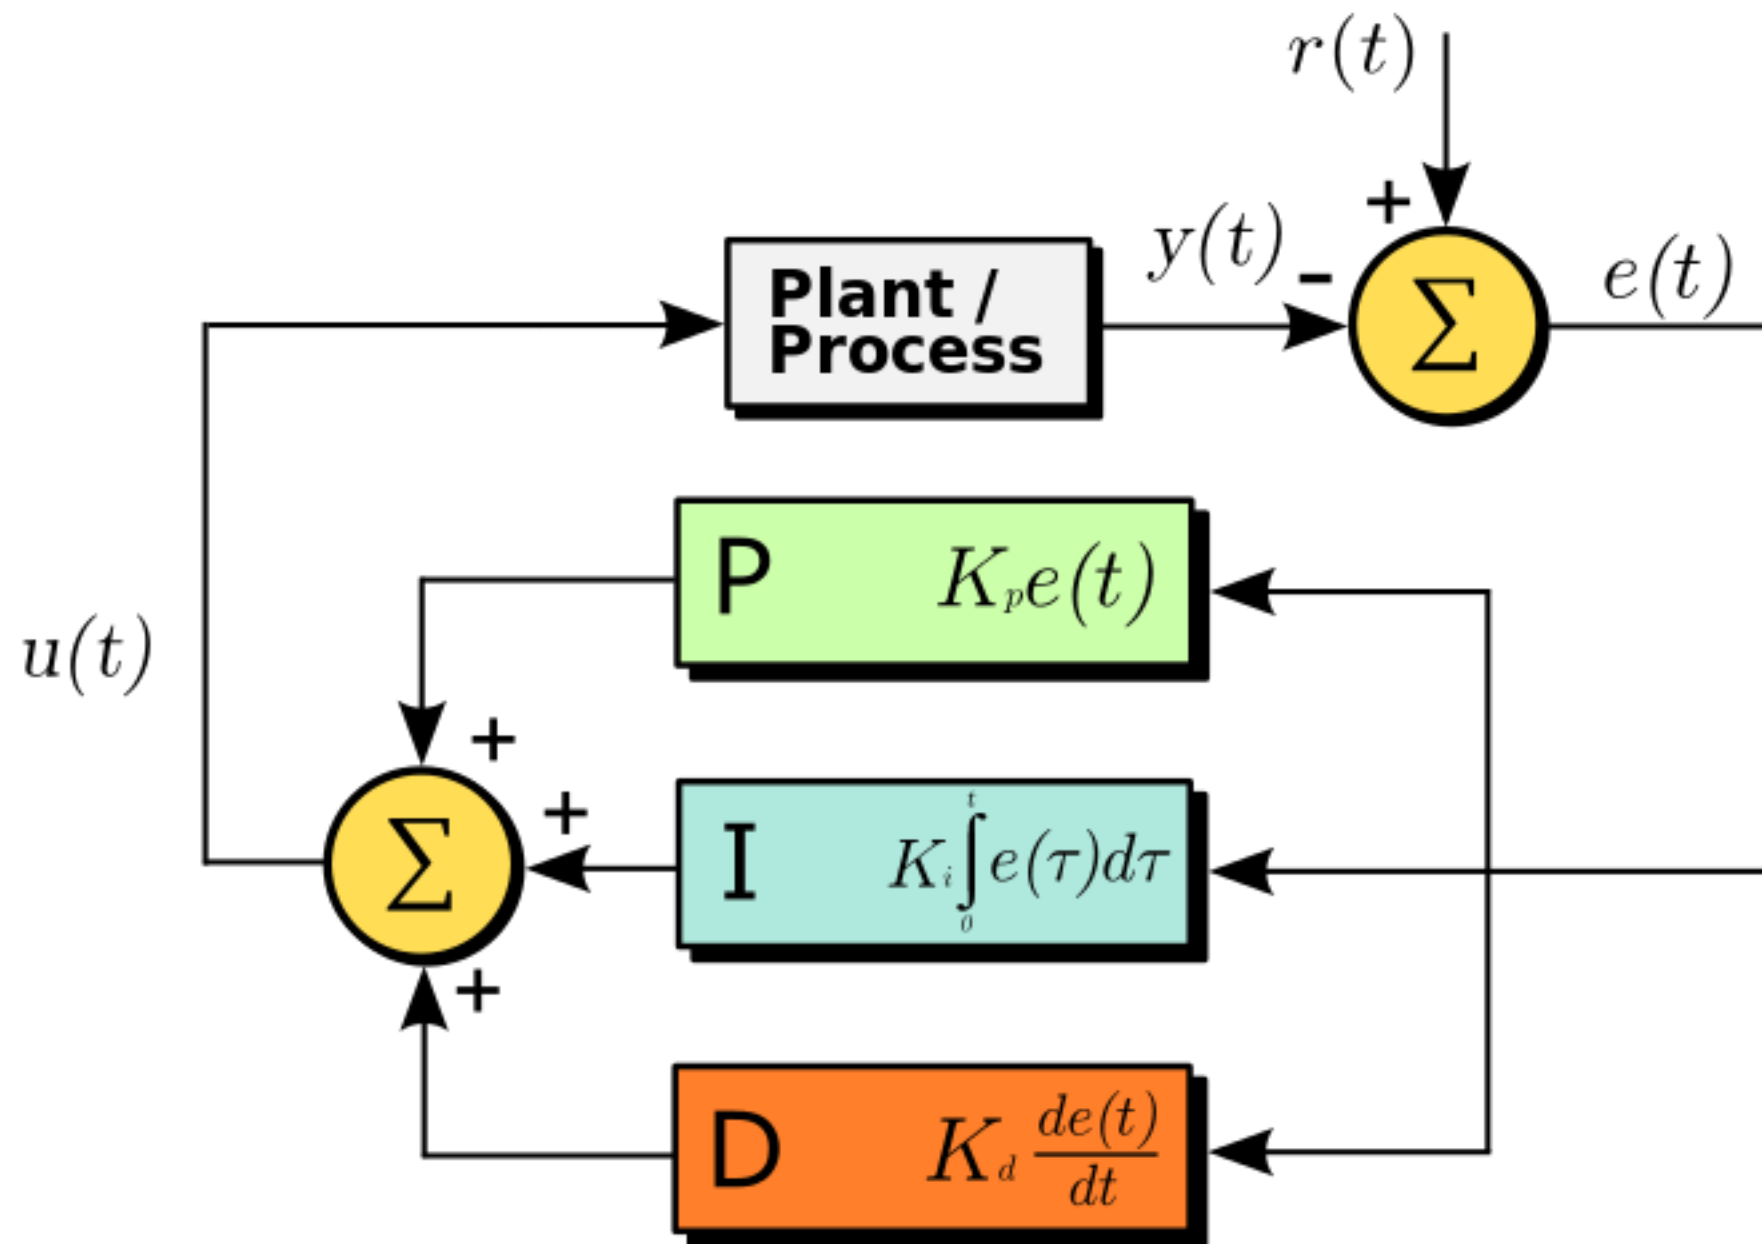
More Details...

- Each `BatchCompleted` event contains
 - processing delay, scheduling delay
 - number of elements in each `InputStream`
- the rate is `elements/processingDelay` (roughly)
 - but what about accumulated data?

PID Controller

- *Proportional, Integrative, Derivative Controller*
- https://en.wikipedia.org/wiki/PID_controller

PID Controller



PID Controller

- Goal: Keep the processing time close to batch interval
- Proportional term: Use the error:
 - $\text{batchInterval} - \text{processingTime}/\text{batch}$
- Integral component: schedulingDelay
- Derived component: $\text{error} - \text{previousError}$

PID Controller

- Not all 3 terms required.
- Does not require process knowledge (black box)
- P, I, D constants can change
convergence, over-shooting, oscillations

To see how the constants affect convergence:

[https://upload.wikimedia.org/wikipedia/commons/3/33/
PID_Compensation_Animated.gif](https://upload.wikimedia.org/wikipedia/commons/3/33/PID_Compensation_Animated.gif)


Back pressure in Spark 1.5

- Each input stream gets its own estimator
- Works with Receiver based inputs, and also `KafkaDirectInputStream`
- To enable, set the following properties:

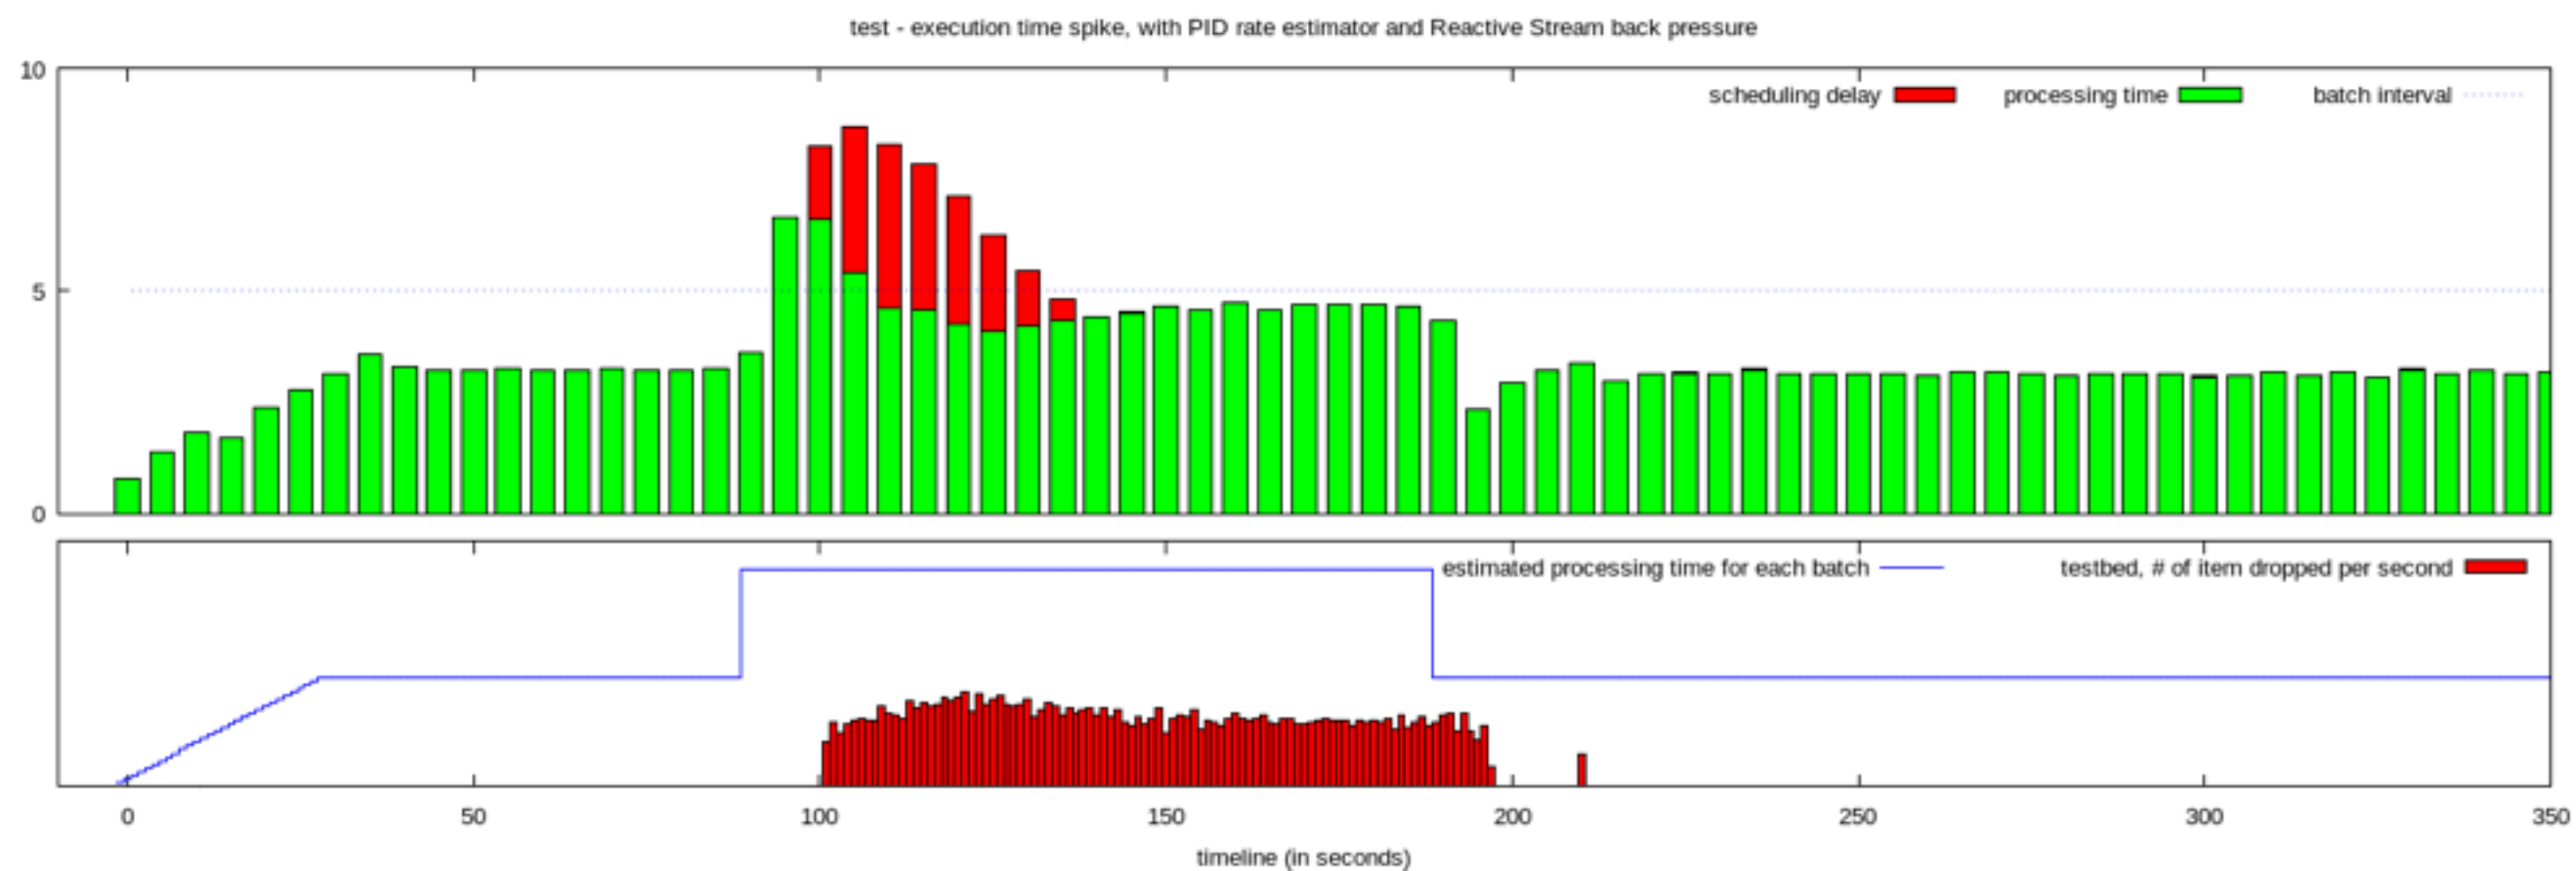

<code>spark.streaming.backpressure.enabled</code>	<code>true</code>
<code>spark.streaming.backpressure.minRate</code>	<code>R</code>

What it does in practice.

Before...



after



Limitations

- Doesn't account for records of different sizes
- Linearity assumption not accurate
- Back pressure accumulates data in the TCP channel; what happens to the data congesting at the source end??

A Peek at Reactive Streams

Reactive Streams

- reactive-streams.org
- A standard for asynchronous stream processing
- Consumer controls rate by asking for elements from producers

Reactive Streams

- A push model is used when the consumer can keep up.
- Switches to a pull model when rate limiting required.
- Dynamically switches back and forth...

Reactive Streams... in Spark?

- Hoped to be ready for v1.6, but TBD:
- JIRA 10420 (Reactive Streams Receiver)

Credits

- Typesafe: Luc Bourlier, Iulian Dragos, Nilanjan Raychaudhuri, Dean Wampler
- (formerly Typesafe) François Garillot
- Databricks: Tathagata Das (TD)