

MongoDB internals

Table of Contents

| | |
|-------------------------------------|---|
| Patterns used..... | 1 |
| Binaries which are built..... | 1 |
| Major classes..... | 1 |
| Server (mongod) execution flow..... | 3 |
| Directory Overview..... | 3 |

Patterns used

There are classes which function as

1. Factory (used to create new instances of a class)
2. Builder (used to mutate an instance)
3. Impl (multiple implementations of the same class)
4. Interface (abstract trait)
5. Listener (to listen to changes on an instance)

There are some singleton objects in every binary. Search for MONGO_INITIALIZER to find them. See calls to `mongo::runGlobalInitializersOrDie()` at start of each executable.

Binaries which are built

`mongod` (server) : from `db/db.cpp`

`mongos` (sharding proxy) : from `s/server.cpp`

`mongo` (client shell) : `shell/dbshell.cpp`

Functions which are shared by `mongos` and `mongod` are placed in library “`db/mongodandmongos`”

Code specific to `mongod` goes into “`serverOnlyFiles`” in the ‘`Sconscript`’ file.

Major classes

DbMessage/DbResponse encapsulate the wire protocol between mongo clients and server.

ServiceContext : Represents the context in which binary is functioning (e.g. mongo sharding proxy or mongo server). “`ServerContextMongod`” is the singleton instance on the “`mongod`” server.

The singleton is retrieved using `getGlobalServiceContext()`. It owns one or more clients.

ClientBasic : From it are derived `Client` (used in server) and `ClientInfo` (used in sharding proxy) classes. `Client` object in the server binary represents a client connection. It has at most one `OperationContext`. See calls to `ServiceContext::makeClient()`.

OperationContext is the equivalent of `Transaction` in MongoDB. Every `OperationContext` in a server with `CurOp` support has a stack of `CurOp` objects. The entry at the top of the stack is used to record timing and resource statistics for the executing operation or suboperation.

RecordCursor : Each `getmore` on a cursor is a separate `OperationContext`. Storage engines only need to implement the derived `SeekableRecordCursor`

WriteUnitOfWork

WriteConcern

RecoveryUnit

CursorManager : singleton

Command : represents a command executed by the proxy or mongo server. All derived commands reside in “commands” dir.

SnapshotManager

Database has `n` Collections

Each **Collection** has `IndexCatalog`, `CollectionCatalogEntry`, `RecordStore`

IndexCatalog has `IndexAccessMethod`

IndexAccessMethod points to `SortedDataInterface`

StorageEngine <- `KVStorageEngine`, `MMapV1Engine`

KVStorageEngine has `KVEngine` as member

KVEngine <- `WiredTiger`, `RocksDB`

KVCatalog <- `RecordStore`

Server (mongod) execution flow

MessageServer::run()

MessageHandler

Request::process

execCommand

parseQuery : convert raw string -> CanonicalQuery -> QuerySolution tree -> PlanStage tree.

PlanExecutor : execute PlanStages against RecordStore and IndexAccessMethod classes.

Directory Overview

base

bson

crypto

client

mongo shell code (clientdriver library)

executor

AsyncStreamFactoryInterface, AsyncStreamInterface,
AsyncTimerFactoryInterface, AsyncTimerInterface,
NetworkInterface
TaskExecutor
ThreadPoolInterface
ConnectionPool

logger

platform

rpc

s - sharding related code

scripting

stdx – related to standard C++ library classes

util

util/concurrency

Locks and threadpool

util/net

- MessageHandler
- MessageServer
- View
- Message

db

- ServiceContext
- CurOp
- OperationContext

db/auth

db/catalog holds engine-independent code to represent/manipulate a column family or index. Calls code in db/storage.

- Collection
- IndexCatalog

db/commands (in turn calls db/query)
 executes commands received from the client
 calls getExecutor() and its variants.

db/concurrency holds lock manager
 LockManager - singleton

db/exec contains code for various PlanStages.

PlanStage (and its 34 derived classes) : This represents a tree of data access and data transforms needed to satisfy a command.

WorkingSet – an operation is executed in many stages. All stages share the working set (i.e. data on which the stages operate).

db/ftdc – stands for full time diagnostic data capture. It takes a set of BSON documents containing metrics, and compresses them into a highly compressed buffers.

db/fts – code to implement text search

db/geo - code for geo indexing

db/index

IndexAccessMethod and derived classes, which call on engine-specific SortedDataInterface implementations.

- IndexDescriptor

db/matcher - compares json with pattern

- MatchExpression and derived classes for different expressions.
- MatchableDocument
- Matcher

db/modules – is symbolically linked to other storage engine

db/ops

db/pipeline holds code for query execution

Expression - derived classes

Pipeline - is used in mongos and mongod

PipelineD is used in mongod

db/query (calls db/exec and db/catalog) : holds code to parse a query string and create query plan

CanonicalQuery : any query is transformed into a canonical representation

PlanExecutor : iterates over tree of PlanStages.

QuerySolution : holds tree of QuerySolutionNodes.

getExecutorUpdate/getExecutorDelete/getExecutorFind all execute in two major steps

-> CanonicalQuery::canonicalize

-> getExecutor

CanonicalQuery::canonicalize : convert raw BSON string -> CanonicalQuery

getExecutor : convert CanonicalQuery -> PlanExecutor

-> prepareForExecution : convert CanonicalQuery->QuerySolution

-> StageBuilder::build : convert QuerySolution -> PlanStage tree

-> PlanExecutor::make : convert CanonicalQuery -> PlanExecutor

LiteParsedQuery

convert BSONObj -> LiteParsedQuery

QueryPlanner::plan

CanonicalQuery -> QuerySolution

PlanExecutor::executePlan

db/repl

DatabaseCloner

CollectionCloner

QuorumChecker

db/s (sharding related)

db/sorter

db/stats

db/storage (storage engine-specific code)

StorageEngine

RecordStore

SnapshotManager